

# Predicting Impact Of Climate Change By Analyzing Climate Data

## Group No: 2

```
In [1]: # Importing Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import warnings
import plotly.express as px
warnings.filterwarnings('ignore')
```

```
In [2]: # Reading the csv file
df_test = pd.read_csv('conventional_weather_stations_inmet_brazil_1961_2019.csv')

df_test.head(5)
```

Out[2]:

	Estacao	Data	Hora	Precipitacao	TempBulboSeco	TempBulboUmido	TempMaxima	Temp
0	82024	01/01/1961	0	NaN	NaN	NaN	32.3	
1	82024	01/01/1961	1200	NaN	26.0	23.9	NaN	
2	82024	01/01/1961	1800	NaN	32.3	27.0	NaN	
3	82024	02/01/1961	0	NaN	25.8	24.6	33.2	
4	82024	02/01/1961	1200	16.0	26.8	24.0	NaN	

```
In [3]: # Taking only the necessary Attributes
# #choosing the necessary attributes to perform the required task on it

df = df_test[['Data', 'Precipitacao', 'TempMaxima', 'TempMinima', 'UmidadeRelativa',
               'DirecaoVento', 'VelocidadeVento', 'Nebulosidade']]
df.head(10)
```

Out[3]:

	Data	Precipitacao	TempMaxima	TempMinima	UmidadeRelativa	DirecaoVento	VelocidadeVento
0	01/01/1961	NaN	32.3	NaN	NaN	NaN	
1	01/01/1961	NaN	NaN	22.9	83.0	5.0	
2	01/01/1961	NaN	NaN	NaN	65.0	5.0	
3	02/01/1961	NaN	33.2	NaN	91.0	9.0	
4	02/01/1961	16.0	NaN	23.7	78.0	5.0	
5	02/01/1961	NaN	NaN	NaN	64.0	5.0	
6	03/01/1961	NaN	32.9	NaN	84.0	5.0	
7	03/01/1961	0.0	NaN	23.3	76.0	9.0	
8	03/01/1961	NaN	NaN	NaN	57.0	5.0	
9	04/01/1961	NaN	30.4	NaN	75.0	5.0	

choosing the necessary attributes to perform the required task on it

```

In [4]: #Renaming Columns
#renaming the columns such that it is easy for us to understand the data which are

df.rename(columns = {
    'Data': 'Date',
    'Precipitacao': 'Precipitation',
    'TempMaxima': 'Maximum_Temperature',
    'TempMinima': 'Minimum_Temperature',
    'UmidadeRelativa': 'Relative_humidity',
    'DirecaoVento': 'Wind_direction',
    'VelocidadeVento': 'Wind_speed',
    'Nebulosidade': 'Cloudiness'
}, inplace = True)
data = df.copy()
df.head(10)

```

Out[4]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Wind_
0	01/01/1961	NaN	32.3	NaN	NaN	
1	01/01/1961	NaN	NaN	22.9	83.0	
2	01/01/1961	NaN	NaN	NaN	65.0	
3	02/01/1961	NaN	33.2	NaN	91.0	
4	02/01/1961	16.0	NaN	23.7	78.0	
5	02/01/1961	NaN	NaN	NaN	64.0	
6	03/01/1961	NaN	32.9	NaN	84.0	
7	03/01/1961	0.0	NaN	23.3	76.0	
8	03/01/1961	NaN	NaN	NaN	57.0	
9	04/01/1961	NaN	30.4	NaN	75.0	

Renaming the columns such that it is easy for us to understand the data which are present in the columns for further work

In [5]: *#Understanding information about the data*  
*#presenting the information regarding the data such that required task can be per*  
 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12251335 entries, 0 to 12251334
Data columns (total 8 columns):
#   Column                Dtype
---  -
0   Date                  object
1   Precipitation          float64
2   Maximum_Temperature    float64
3   Minimum_Temperature     float64
4   Relative_humidity       float64
5   Wind_direction          float64
6   Wind_speed             float64
7   Cloudiness             float64
dtypes: float64(7), object(1)
memory usage: 747.8+ MB
```

presenting the information regarding the data such that required task can be performed on it

In [6]: *#Checking for missing values*

```
print(df.isnull().sum())
```

```
Date                0
Precipitation        8130580
Maximum_Temperature  8300413
Minimum_Temperature  8292199
Relative_humidity     846548
Wind_direction       1221386
Wind_speed           1019529
Cloudiness           433425
dtype: int64
```

checking whether there are missing values to understand about noisy data

In [7]: *#checking missing columns in Precipitation*

```
df[df["Precipitation"].isnull()]
```

Out[7]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity
0	01/01/1961	NaN	32.3	NaN	NaN
1	01/01/1961	NaN	NaN	22.9	83.0
2	01/01/1961	NaN	NaN	NaN	65.0
3	02/01/1961	NaN	33.2	NaN	91.0
5	02/01/1961	NaN	NaN	NaN	64.0
...	...	...	...	...	..
12251328	29/12/2019	NaN	NaN	NaN	43.0
12251329	30/12/2019	NaN	32.1	NaN	60.0
12251331	30/12/2019	NaN	NaN	NaN	46.0
12251332	31/12/2019	NaN	31.6	NaN	55.0
12251334	31/12/2019	NaN	NaN	NaN	43.0

8130580 rows × 8 columns

In [8]: *# filling the missing values with their mean*

```
#
meanVal1 = df['Precipitation'].mean()
df['Precipitation'].fillna(value=meanVal1, inplace=True)

meanVal2 = df['Maximum_Temperature'].mean()
df['Maximum_Temperature'].fillna(value=meanVal2, inplace=True)

meanVal3 = df['Minimum_Temperature'].mean()
df['Minimum_Temperature'].fillna(value=meanVal3, inplace=True)

meanVal4 = df['Relative_humidity'].mean()
df['Relative_humidity'].fillna(value=meanVal4, inplace=True)

meanVal5 = df['Wind_direction'].mean()
df['Wind_direction'].fillna(value=meanVal5, inplace=True)

meanVal6 = df['Wind_speed'].mean()
df['Wind_speed'].fillna(value=meanVal6, inplace=True)

meanVal7 = df['Cloudiness'].mean()
df['Cloudiness'].fillna(value=meanVal7, inplace=True)
```

There are certain methods which can be used to replace missing values we here is using the mean method such that the computation is easier

In [9]: *# checking for missing values*

```
df.isnull().sum()
```

```
Out[9]: Date                0
Precipitation              0
Maximum_Temperature        0
Minimum_Temperature        0
Relative_humidity          0
Wind_direction             0
Wind_speed                 0
Cloudiness                 0
dtype: int64
```

In [10]: `df.head(10)`

Out[10]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Wind_
0	01/01/1961	4.027972	32.300000	18.927569	72.21245	
1	01/01/1961	4.027972	29.725229	22.900000	83.00000	
2	01/01/1961	4.027972	29.725229	18.927569	65.00000	
3	02/01/1961	4.027972	33.200000	18.927569	91.00000	
4	02/01/1961	16.000000	29.725229	23.700000	78.00000	
5	02/01/1961	4.027972	29.725229	18.927569	64.00000	
6	03/01/1961	4.027972	32.900000	18.927569	84.00000	
7	03/01/1961	0.000000	29.725229	23.300000	76.00000	
8	03/01/1961	4.027972	29.725229	18.927569	57.00000	
9	04/01/1961	4.027972	30.400000	18.927569	75.00000	



In [11]: *#Checking the number of rows and columns*

```
df.shape
```

Out[11]: (12251335, 8)

In [12]: *#Printing the last 5 rows*

```
df.tail()
```

Out[12]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity
<b>12251330</b>	30/12/2019	0.000000	29.725229	21.100000	66.0
<b>12251331</b>	30/12/2019	4.027972	29.725229	18.927569	46.0
<b>12251332</b>	31/12/2019	4.027972	31.600000	18.927569	55.0
<b>12251333</b>	31/12/2019	0.000000	29.725229	21.300000	54.0
<b>12251334</b>	31/12/2019	4.027972	29.725229	18.927569	43.0

In [13]: `df['Date']=pd.to_datetime(df['Date'])`  
`df.head(5)`

Out[13]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Wind_direction
<b>0</b>	1961-01-01	4.027972	32.300000	18.927569	72.21245	12.7
<b>1</b>	1961-01-01	4.027972	29.725229	22.900000	83.00000	5.0
<b>2</b>	1961-01-01	4.027972	29.725229	18.927569	65.00000	5.0
<b>3</b>	1961-02-01	4.027972	33.200000	18.927569	91.00000	9.0
<b>4</b>	1961-02-01	16.000000	29.725229	23.700000	78.00000	5.0

In [14]: *#Descriptive Statistics*  
*#*

```
df.describe()
```

Out[14]:

	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Wind_direction
<b>count</b>	1.225134e+07	1.225134e+07	1.225134e+07	1.225134e+07	1.225134e+
<b>mean</b>	4.027972e+00	2.972523e+01	1.892757e+01	7.221245e+01	1.274881e+
<b>std</b>	6.321910e+00	2.592241e+00	2.528762e+00	1.746268e+01	1.053072e+
<b>min</b>	0.000000e+00	-2.000000e+00	-9.000000e+00	8.000000e+00	0.000000e+
<b>25%</b>	1.600000e+00	2.972523e+01	1.892757e+01	6.100000e+01	5.000000e+
<b>50%</b>	4.027972e+00	2.972523e+01	1.892757e+01	7.300000e+01	1.274881e+
<b>75%</b>	4.027972e+00	2.972523e+01	1.892757e+01	8.600000e+01	1.800000e+
<b>max</b>	3.779000e+02	4.470000e+01	3.650000e+01	1.000000e+02	9.900000e+

## Describing the required columns

```
In [15]: #Normalization Using Min Max Scaling
#

x = df.iloc[:,1:7]
global_min = x.min()
global_max = x.max()
df.iloc[:,1:7] = (x-global_min) / (global_max - global_min)

df.describe()
```

Out[15]:

	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Wind_direction
<b>count</b>	1.225134e+07	1.225134e+07	1.225134e+07	1.225134e+07	1.225134e+07
<b>mean</b>	1.065883e-02	6.793411e-01	6.137927e-01	6.979614e-01	1.287759e-01
<b>std</b>	1.672906e-02	5.550837e-02	5.557719e-02	1.898117e-01	1.063709e-01
<b>min</b>	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
<b>25%</b>	4.233924e-03	6.793411e-01	6.137927e-01	5.760870e-01	5.050505e-02
<b>50%</b>	1.065883e-02	6.793411e-01	6.137927e-01	7.065217e-01	1.287759e-01
<b>75%</b>	1.065883e-02	6.793411e-01	6.137927e-01	8.478261e-01	1.818182e-01
<b>max</b>	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00

The process of transforming the columns in a dataset to the same scale



In [16]: *# Getting year, month, day and day name*

```
df['Year']=df["Date"].dt.year
df['Month']=df["Date"].dt.month
df['Day']=df["Date"].dt.day
df['Day Name']=df["Date"].dt.day_name()
df.head(10)
```

Out[16]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Wind_direction
0	1961-01-01	0.010659	0.734475	0.613793	0.697961	0.12871
1	1961-01-01	0.010659	0.679341	0.701099	0.815217	0.05050
2	1961-01-01	0.010659	0.679341	0.613793	0.619565	0.05050
3	1961-02-01	0.010659	0.753747	0.613793	0.902174	0.09090
4	1961-02-01	0.042339	0.679341	0.718681	0.760870	0.05050
5	1961-02-01	0.010659	0.679341	0.613793	0.608696	0.05050
6	1961-03-01	0.010659	0.747323	0.613793	0.826087	0.05050
7	1961-03-01	0.000000	0.679341	0.709890	0.739130	0.09090
8	1961-03-01	0.010659	0.679341	0.613793	0.532609	0.05050
9	1961-04-01	0.010659	0.693790	0.613793	0.728261	0.05050

In [17]: *# Sorting the data in ascending order on the basis of date*

```
df.sort_values(by="Date",ascending=True,inplace=True)
df.head(10)
```

Out[17]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Win
0	1961-01-01	0.010659	0.734475	0.613793	0.697961	
4227197	1961-01-01	0.010659	0.679341	0.613793	0.293478	
4227196	1961-01-01	0.010659	0.679341	0.727473	0.489130	
4227195	1961-01-01	0.010659	0.824411	0.613793	0.697961	
3206034	1961-01-01	0.010659	0.608137	0.613793	0.697961	
3206035	1961-01-01	0.010659	0.679341	0.624176	0.782609	
3206036	1961-01-01	0.010659	0.679341	0.613793	0.521739	
7519668	1961-01-01	0.010659	0.725910	0.613793	0.697961	
7519669	1961-01-01	0.010659	0.679341	0.595604	0.771739	
7519670	1961-01-01	0.010659	0.679341	0.613793	0.597826	



```
In [18]: df.tail(10)
```

Out[18]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Wi
<b>2421854</b>	2019-12-31	0.010659	0.679341	0.613793	0.456522	
<b>6253606</b>	2019-12-31	0.010659	0.679341	0.613793	0.467391	
<b>6253605</b>	2019-12-31	0.000000	0.679341	0.690110	0.565217	
<b>6253604</b>	2019-12-31	0.010659	0.764454	0.613793	0.673913	
<b>6209680</b>	2019-12-31	0.010659	0.679341	0.613793	0.608696	
<b>6209679</b>	2019-12-31	0.000000	0.679341	0.661538	0.706522	
<b>6209678</b>	2019-12-31	0.010659	0.732334	0.613793	0.673913	
<b>6166564</b>	2019-12-31	0.010659	0.679341	0.613793	0.478261	
<b>12195887</b>	2019-12-31	0.010659	0.679341	0.613793	0.923913	
<b>12251334</b>	2019-12-31	0.010659	0.679341	0.613793	0.380435	

```
In [19]: df.head(10)
```

Out[19]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Win
0	1961-01-01	0.010659	0.734475	0.613793	0.697961	
4227197	1961-01-01	0.010659	0.679341	0.613793	0.293478	
4227196	1961-01-01	0.010659	0.679341	0.727473	0.489130	
4227195	1961-01-01	0.010659	0.824411	0.613793	0.697961	
3206034	1961-01-01	0.010659	0.608137	0.613793	0.697961	
3206035	1961-01-01	0.010659	0.679341	0.624176	0.782609	
3206036	1961-01-01	0.010659	0.679341	0.613793	0.521739	
7519668	1961-01-01	0.010659	0.725910	0.613793	0.697961	
7519669	1961-01-01	0.010659	0.679341	0.595604	0.771739	
7519670	1961-01-01	0.010659	0.679341	0.613793	0.597826	

In [20]: *# Creating a dataframe for the year 1961 year records*

```
df_2=df[(df['Year'] == 1961)]
df_2.head(10)
```

Out[20]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Win
0	1961-01-01	0.010659	0.734475	0.613793	0.697961	
4227197	1961-01-01	0.010659	0.679341	0.613793	0.293478	
4227196	1961-01-01	0.010659	0.679341	0.727473	0.489130	
4227195	1961-01-01	0.010659	0.824411	0.613793	0.697961	
3206034	1961-01-01	0.010659	0.608137	0.613793	0.697961	
3206035	1961-01-01	0.010659	0.679341	0.624176	0.782609	
3206036	1961-01-01	0.010659	0.679341	0.613793	0.521739	
7519668	1961-01-01	0.010659	0.725910	0.613793	0.697961	
7519669	1961-01-01	0.010659	0.679341	0.595604	0.771739	
7519670	1961-01-01	0.010659	0.679341	0.613793	0.597826	



In [21]: *# Grouping the records for 1961 dataframe by month*

```
group_1961_month=df_2.groupby('Month')
```

In [22]:

```
# Getting the records for the month of November
df_2[df_2['Month'] == 11]
```

Out[22]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Win
7731764	1961-11-01	0.000000	0.679341	0.509890	0.456522	
6670104	1961-11-01	0.010659	0.679341	0.613793	0.706522	
7731763	1961-11-01	0.010659	0.695931	0.613793	0.554348	
7731765	1961-11-01	0.010659	0.679341	0.613793	0.315217	
6670103	1961-11-01	0.041810	0.679341	0.610989	0.880435	
...	...	...	...	...	...	...
3668799	1961-11-30	0.000529	0.679341	0.597802	0.673913	
3668798	1961-11-30	0.010659	0.770878	0.613793	0.717391	
1523885	1961-11-30	0.010659	0.809422	0.613793	0.489130	
3150393	1961-11-30	0.010659	0.679341	0.613793	0.402174	
3206946	1961-11-30	0.010659	0.679341	0.613793	0.391304	

11526 rows × 12 columns



In [23]:

# Getting the records for the month of December  
df\_2[df\_2['Month'] == 12]

Out[23]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Wi
12195922	1961-12-01	0.069331	0.679341	0.567033	0.945652	
12195921	1961-12-01	0.010659	0.556745	0.613793	0.967391	
3316000	1961-12-01	0.010659	0.616702	0.613793	0.869565	
4273221	1961-12-01	0.010659	0.679341	0.613793	0.586957	
9278724	1961-12-01	0.023816	0.679341	0.624176	0.923913	
...	...	...	...	...	...	...
8562199	1961-12-31	0.010659	0.679341	0.613793	0.695652	
87079	1961-12-31	0.010659	0.672377	0.613793	0.978261	
7225322	1961-12-31	0.000000	0.679341	0.490110	0.663043	
1476254	1961-12-31	0.000000	0.679341	0.571429	0.815217	
1476253	1961-12-31	0.010659	0.642398	0.613793	0.760870	

11325 rows × 12 columns



In [24]: *# Creating a dataframe for the year 2019 year records*

```
df_3=df[(df['Year'] == 2019)]
df_3.head(10)
```

Out[24]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Win
<b>2421013</b>	2019-01-01	0.010659	0.719486	0.613793	0.956522	
<b>2421014</b>	2019-01-01	0.002911	0.679341	0.641758	0.815217	
<b>9182776</b>	2019-01-01	0.010659	0.679341	0.613793	0.293478	
<b>9182775</b>	2019-01-01	0.000000	0.679341	0.716484	0.543478	
<b>8129162</b>	2019-01-01	0.010659	0.728051	0.613793	0.826087	
<b>2987037</b>	2019-01-01	0.010659	0.670236	0.613793	0.836957	
<b>1856439</b>	2019-01-01	0.108759	0.679341	0.685714	0.804348	
<b>1856438</b>	2019-01-01	0.010659	0.721627	0.613793	0.967391	
<b>143531</b>	2019-01-01	0.010659	0.679341	0.613793	0.663043	
<b>143532</b>	2019-01-01	0.000000	0.679341	0.613793	0.945652	



In [25]: *# Grouping the records for 2019 dataframe by month*

```
group_2019_month=df_2.groupby('Month')
```



In [26]: *# Getting the records for the month of November 2019*

```
df_3[df_3['Month'] == 11]
```

Out[26]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Wi
<b>9624162</b>	2019-11-01	0.010659	0.679341	0.613793	0.532609	
<b>1572690</b>	2019-11-01	0.010659	0.679341	0.613793	0.697961	
<b>10447867</b>	2019-11-01	0.010659	0.679341	0.613793	0.163043	
<b>2473333</b>	2019-11-01	0.010659	0.704497	0.613793	0.782609	
<b>7673930</b>	2019-11-01	0.010659	0.679341	0.613793	0.750000	
...	...	...	...	...	...	...
<b>1307171</b>	2019-11-30	0.010659	0.753747	0.613793	0.608696	
<b>10525892</b>	2019-11-30	0.000000	0.679341	0.654945	0.641304	
<b>10525893</b>	2019-11-30	0.010659	0.679341	0.613793	0.315217	
<b>10525891</b>	2019-11-30	0.010659	0.679341	0.613793	0.630435	
<b>10306388</b>	2019-11-30	0.010659	0.679341	0.613793	0.697961	

15256 rows × 12 columns



In [27]: *# Descriptive statistics for November 1961*

```
df_2[df_2['Month'] == 11].describe()
```

Out[27]:

	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Wind_direction
<b>count</b>	11526.000000	11526.000000	11526.000000	11526.000000	11526.000000
<b>mean</b>	0.010268	0.681382	0.611339	0.679285	0.116300
<b>std</b>	0.014870	0.052272	0.047805	0.187579	0.111100
<b>min</b>	0.000000	0.226981	0.215385	0.065217	0.000000
<b>25%</b>	0.004234	0.679341	0.613793	0.565217	0.000000
<b>50%</b>	0.010659	0.679341	0.613793	0.697961	0.090900
<b>75%</b>	0.010659	0.679341	0.613793	0.826087	0.181800
<b>max</b>	0.339772	0.905782	0.795604	1.000000	0.363600



In [28]: *# Descriptive statistics for November 2005*

```
df_3[df_3['Month'] == 11].describe()
```

Out[28]:

	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Wind_direction
<b>count</b>	15256.000000	15256.000000	15256.000000	15256.000000	15256.000000
<b>mean</b>	0.010914	0.693431	0.628259	0.680992	0.125500
<b>std</b>	0.019225	0.051442	0.049350	0.190286	0.102600
<b>min</b>	0.000000	0.370450	0.237363	0.054348	0.000000
<b>25%</b>	0.001058	0.679341	0.613793	0.554348	0.050500
<b>50%</b>	0.010659	0.679341	0.613793	0.697961	0.128700
<b>75%</b>	0.010659	0.679341	0.619780	0.826087	0.141400
<b>max</b>	0.464144	0.942184	0.848352	1.000000	0.363600



In [29]:

# Descriptive Statistics for all the month for the year 1961

df\_2.groupby('Month').describe()

Out[29]:

Month	Precipitation								Maximum_Temperature	
	count	mean	std	min	25%	50%	75%	max	count	mean
1	12028.0	0.012500	0.019973	0.0	0.010659	0.010659	0.010659	0.361736	12028.0	0.6
2	10795.0	0.012589	0.021801	0.0	0.010659	0.010659	0.010659	0.429214	10795.0	0.6
3	11936.0	0.011657	0.019052	0.0	0.010659	0.010659	0.010659	0.358296	11936.0	0.6
4	11643.0	0.011157	0.018138	0.0	0.007674	0.010659	0.010659	0.420217	11643.0	0.6
5	11964.0	0.010214	0.015903	0.0	0.002646	0.010659	0.010659	0.351680	11964.0	0.6
6	11606.0	0.010086	0.014774	0.0	0.003771	0.010659	0.010659	0.288436	11606.0	0.6
7	11458.0	0.009441	0.012018	0.0	0.001852	0.010659	0.010659	0.237894	11458.0	0.6
8	11557.0	0.009181	0.011527	0.0	0.000529	0.010659	0.010659	0.216459	11557.0	0.6
9	11668.0	0.009620	0.013894	0.0	0.000265	0.010659	0.010659	0.313840	11668.0	0.6
10	12026.0	0.010561	0.016263	0.0	0.005557	0.010659	0.010659	0.354856	12026.0	0.6
11	11526.0	0.010268	0.014870	0.0	0.004234	0.010659	0.010659	0.339772	11526.0	0.6
12	11325.0	0.011476	0.017250	0.0	0.010585	0.010659	0.010659	0.341360	11325.0	0.6

12 rows × 72 columns



In [30]:

# Descriptive Statistics for all the month for the year 2019

df\_3.groupby('Month').describe()

Out[30]:

	Precipitation								Maximum_Temperature	
	count	mean	std	min	25%	50%	75%	max	count	mean
Month										
1	16023.0	0.010881	0.016829	0.0	0.002382	0.010659	0.010659	0.429743	16023.0	0.6
2	14964.0	0.012337	0.020475	0.0	0.008203	0.010659	0.010659	0.343212	14964.0	0.6
3	16247.0	0.012226	0.021521	0.0	0.007145	0.010659	0.010659	0.620270	16247.0	0.6
4	15811.0	0.011087	0.017632	0.0	0.003175	0.010659	0.010659	0.375496	15811.0	0.6
5	16350.0	0.010698	0.017785	0.0	0.001588	0.010659	0.010659	0.525536	16350.0	0.6
6	15249.0	0.009633	0.014349	0.0	0.000265	0.010659	0.010659	0.482667	15249.0	0.6
7	15170.0	0.009709	0.014335	0.0	0.000000	0.010659	0.010659	0.347976	15170.0	0.6
8	16013.0	0.008976	0.012982	0.0	0.000000	0.010659	0.010659	0.451442	16013.0	0.6
9	15131.0	0.009199	0.013176	0.0	0.000000	0.010659	0.010659	0.378936	15131.0	0.6
10	15642.0	0.010115	0.015872	0.0	0.000000	0.010659	0.010659	0.379995	15642.0	0.6
11	15256.0	0.010914	0.019225	0.0	0.001058	0.010659	0.010659	0.464144	15256.0	0.6
12	15213.0	0.010389	0.016873	0.0	0.000529	0.010659	0.010659	0.311723	15213.0	0.6

12 rows × 72 columns



In [31]:


# Getting mean,std deviation and median for all the months in the year 1961

df\_2.groupby('Month').agg(['mean','std','median'])

Out[31]:

	Precipitation			Maximum_Temperature			Minimum_Temperature		
	mean	std	median	mean	std	median	mean	std	median
Month									
1	0.012500	0.019973	0.010659	0.676707	0.045991	0.679341	0.613500	0.048226	0.613793
2	0.012589	0.021801	0.010659	0.675353	0.046444	0.679341	0.612712	0.044793	0.613793
3	0.011657	0.019052	0.010659	0.675365	0.048482	0.679341	0.609511	0.049326	0.613793
4	0.011157	0.018138	0.010659	0.670633	0.054232	0.679341	0.605623	0.055709	0.613793
5	0.010214	0.015903	0.010659	0.665902	0.057372	0.679341	0.597066	0.066128	0.613793
6	0.010086	0.014774	0.010659	0.658584	0.074512	0.679341	0.591006	0.074992	0.613793
7	0.009441	0.012018	0.010659	0.663706	0.063810	0.679341	0.589542	0.074648	0.613793
8	0.009181	0.011527	0.010659	0.673364	0.057476	0.679341	0.597310	0.062103	0.613793
9	0.009620	0.013894	0.010659	0.679690	0.062444	0.679341	0.604803	0.054289	0.613793
10	0.010561	0.016263	0.010659	0.674995	0.055738	0.679341	0.609706	0.049446	0.613793
11	0.010268	0.014870	0.010659	0.681382	0.052272	0.679341	0.611339	0.047805	0.613793
12	0.011476	0.017250	0.010659	0.676231	0.049950	0.679341	0.612787	0.046035	0.613793

12 rows × 27 columns



In [32]: *# Getting mean,std deviation and median for all the months in the year 2019*

```
df_3.groupby('Month').agg(['mean','std','median'])
```

Out[32]:

Month	Precipitation			Maximum_Temperature			Minimum_Temperature		
	mean	std	median	mean	std	median	mean	std	median
1	0.010881	0.016829	0.010659	0.696400	0.047254	0.679341	0.629658	0.044716	0.613793
2	0.012337	0.020475	0.010659	0.691441	0.047328	0.679341	0.629582	0.047836	0.613793
3	0.012226	0.021521	0.010659	0.689203	0.048305	0.679341	0.628573	0.048885	0.613793
4	0.011087	0.017632	0.010659	0.688245	0.048775	0.679341	0.626915	0.049804	0.613793
5	0.010698	0.017785	0.010659	0.682439	0.056432	0.679341	0.620809	0.053980	0.613793
6	0.009633	0.014349	0.010659	0.681463	0.052252	0.679341	0.614618	0.056774	0.613793
7	0.009709	0.014335	0.010659	0.678874	0.058404	0.679341	0.607955	0.063788	0.613793
8	0.008976	0.012982	0.010659	0.684877	0.060634	0.679341	0.612358	0.061553	0.613793
9	0.009199	0.013176	0.010659	0.693958	0.062611	0.679341	0.621193	0.054348	0.613793
10	0.010115	0.015872	0.010659	0.694739	0.055981	0.679341	0.625494	0.049311	0.613793
11	0.010914	0.019225	0.010659	0.693431	0.051442	0.679341	0.628259	0.049350	0.613793
12	0.010389	0.016873	0.010659	0.696196	0.050021	0.679341	0.630388	0.048678	0.613793

12 rows × 27 columns

In [33]: *# Getting mean,sd deviation and median values for the winter season of 1961*

```
df_2[(df_2['Month'] == 12) | (df_2['Month'] == 1) | (df_2['Month'] == 2)].groupby
```

Out[33]:

Month	Precipitation			Maximum_Temperature			Minimum_Temperature		
	mean	std	median	mean	std	median	mean	std	median
1	0.012500	0.019973	0.010659	0.676707	0.045991	0.679341	0.613500	0.048226	0.613793
2	0.012589	0.021801	0.010659	0.675353	0.046444	0.679341	0.612712	0.044793	0.613793
12	0.011476	0.017250	0.010659	0.676231	0.049950	0.679341	0.612787	0.046035	0.613793

3 rows × 27 columns

```
In [34]: # Grouping the dataframe for the year 1961 about 'Month' and 'Day Name' columns
df_2.groupby(['Month', 'Day Name']).count()
```

Out[34]:

		Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_hun
Month	Day Name					
1	Friday	1560	1560	1560	1560	
	Monday	1940	1940	1940	1940	
	Saturday	1547	1547	1547	1547	
	Sunday	1919	1919	1919	1919	
	Thursday	1553	1553	1553	1553	
...	...	...	...	...	...	...
12	Saturday	1842	1842	1842	1842	
	Sunday	1846	1846	1846	1846	
	Thursday	1428	1428	1428	1428	
	Tuesday	1450	1450	1450	1450	
	Wednesday	1454	1454	1454	1454	

84 rows × 10 columns



```
In [35]: # Getting the count of each day for each month in 1961
df_2.groupby(['Month', 'Day Name']).count()
```

Out[35]:

		Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_hun
Month	Day Name					
1	Friday	1560	1560	1560	1560	
	Monday	1940	1940	1940	1940	
	Saturday	1547	1547	1547	1547	
	Sunday	1919	1919	1919	1919	
	Thursday	1553	1553	1553	1553	
...	...	...	...	...	...	...
12	Saturday	1842	1842	1842	1842	
	Sunday	1846	1846	1846	1846	
	Thursday	1428	1428	1428	1428	
	Tuesday	1450	1450	1450	1450	
	Wednesday	1454	1454	1454	1454	

84 rows × 10 columns




In [36]: `df.head(10)`

Out[36]:

	Date	Precipitation	Maximum_Temperature	Minimum_Temperature	Relative_humidity	Win
0	1961-01-01	0.010659	0.734475	0.613793	0.697961	
4227197	1961-01-01	0.010659	0.679341	0.613793	0.293478	
4227196	1961-01-01	0.010659	0.679341	0.727473	0.489130	
4227195	1961-01-01	0.010659	0.824411	0.613793	0.697961	
3206034	1961-01-01	0.010659	0.608137	0.613793	0.697961	
3206035	1961-01-01	0.010659	0.679341	0.624176	0.782609	
3206036	1961-01-01	0.010659	0.679341	0.613793	0.521739	
7519668	1961-01-01	0.010659	0.725910	0.613793	0.697961	
7519669	1961-01-01	0.010659	0.679341	0.595604	0.771739	
7519670	1961-01-01	0.010659	0.679341	0.613793	0.597826	

## Data Visualization

In [37]: `# Plotting the variation in precipitation in Brazil between the years 1961 to 2019`  
`#`  
`plt.figure(figsize = (25,5));`  
`plt.plot(data.Precipitation);`  
`plt.ylabel('Precipitation');`  
`plt.xlabel('Years')`  
`plt.title('Variations in Precipitation in Brazilia 1961-2020');`



From the graph there is variance in the precipitation level in Brazil from 1961 to 2019. There is a certain rise and fall of the level.



```
In [38]: # Plotting the yearly maximum temperature and the year


df.index = pd.to_datetime(df['Date'])
average_monthly_temperature = df['Maximum_Temperature'].groupby(df.index.year).me

x = np.arange(1961, 2020)

l = average_monthly_temperature.tolist()

average_monthly_temperature
plt.plot(x, l, label="line L")
```

Out[38]: [<matplotlib.lines.Line2D at 0x20d8587f5e0>]




The above graph shows us the variance of maximum temperature to that of the year and we can see there is a huge rise in temperature in 2020 and there is a downfall in between 1990 to 2000.

```
In [39]: # Drawing a scatterplot between precipitation and cloudiness

col1 = df['Precipitation']
col2 = df['Cloudiness']

plt.figure(figsize = (20, 5))
plt.scatter(col1, col2,color="blue",)
plt.xlabel('Precipitation')
plt.ylabel('Cloudiness')
plt.show()
```



The above graph shows us how precipitation is affected by cloudiness

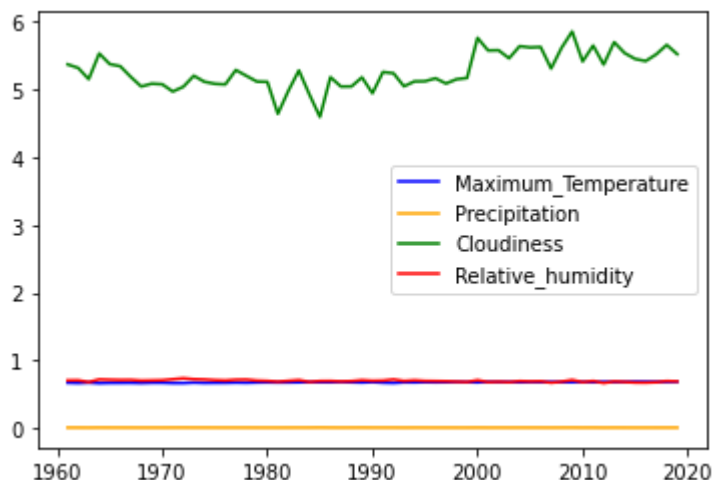
```
In [40]: # Plotting the relationship between cloudiness, temperature, precipitation and relative humidity

average_monthly_temperature = df['Maximum_Temperature'].groupby(df.index.year).mean()
average_monthly_precipitation = df['Precipitation'].groupby(df.index.year).mean()
average_monthly_cloud = df['Cloudiness'].groupby(df.index.year).mean()
average_monthly_humidity = df['Relative_humidity'].groupby(df.index.year).mean()

a = average_monthly_temperature.tolist()
b = average_monthly_precipitation.tolist()
c = average_monthly_cloud.tolist()
d = average_monthly_humidity.tolist()

plt.plot(x, a, label="Maximum_Temperature", color="blue")
plt.plot(x, b, label="Precipitation", color="orange")
plt.plot(x, c, label="Cloudiness", color="green")
plt.plot(x, d, label="Relative_humidity", color="red")
plt.legend()
```

Out[40]: <matplotlib.legend.Legend at 0x20d21467970>




The above graph describes that precipitation, maximum temperature and relative humidity is almost constant but the cloudiness varies so much throughout

```
In [41]: # Boxplot
```

```
df.boxplot(column=['Precipitation','Maximum_Temperature','Relative_humidity','Wind_speed','Cloudiness'])
```

```
Out[41]: <AxesSubplot:>
```




The box plot shows the distribution of the attributes it is being compared with one another.

```
In [42]: # Plotting the histogram for the precipitation
```

```
df['Precipitation'].hist(bins=20,figsize=(20,10))
```

```
Out[42]: <AxesSubplot:>
```




The graph displays the continues values which explains the rise and then fall of precipitation

In [43]: *# Plotting the histogram for the maximum temperature*

```
df['Maximum_Temperature'].hist(bins=20,figsize=(20,10))
```

Out[43]: <AxesSubplot:>



The above graph shows the continues values from low to a certain rise and going low of the attibute Maximum\_temperature

In [44]: *# Create twodimensional NumPy arrays from onedimensional Pandas series for the fe*


```
X = df[['Precipitation', 'Maximum_Temperature', 'Relative_humidity', 'Wind_speed', 'C
y = df['Year'].values
print(X)
print(y)
```

```
[[ 0.01065883  0.73447537  0.69796141  0.08104219  5.30951457]
 [ 0.01065883  0.67934109  0.29347826  0.04135875  5.          ]
 [ 0.01065883  0.67934109  0.48913043  0.0827175   4.          ]
 ...
 [ 0.01065883  0.67934109  0.47826087  0.08510638 10.          ]
 [ 0.01065883  0.67934109  0.92391304  0.04255319 10.          ]
 [ 0.01065883  0.67934109  0.38043478  0.          8.75         ]]
[1961 1961 1961 ... 2019 2019 2019]
```

In [45]: *# Apply the 'norm.pdf()' function to get the probabilities and then create a norm*

```
import scipy.stats as stats
import pylab
```

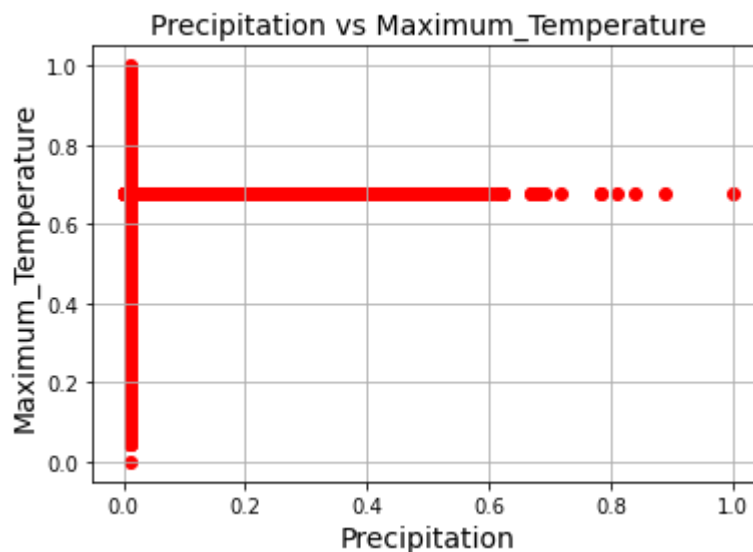
```
stats.probplot(df['Precipitation'], dist="norm", plot=pylab)
pylab.show()
```



This graph displays the probabilities of theoretical quantiles to that of ordered values. Here percipitaion is having a rise in (2,1.0)

In [46]: *# Create a scatter plot between the errors and the independent variable for the t*

```
plt.scatter(df['Precipitation'], df['Maximum_Temperature'], color = 'red')
plt.title('Precipitation vs Maximum_Temperature', fontsize = 14)
plt.xlabel('Precipitation', fontsize = 14)
plt.ylabel('Maximum_Temperature', fontsize = 14)
plt.grid(True)
plt.show()
```




The graph between Maximum\_temperature to that of percipitation.Here the graph displays how much Temperature is affected with percipitation

In [47]: *#Time-Series Line Plots*

```
df['Precipitation'].plot(figsize=(20,10))
```

Out[47]: <AxesSubplot:xlabel='Date'>




This graph is a time graph it tells us that during 1990 the percipitation is the highest recorded according to the data.It is a very much scattered garph

In [48]: *# Time-Series Line Plots*

```
df['Maximum_Temperature'].plot(figsize=(20,10))
```

Out[48]: <AxesSubplot:xlabel='Date'>




This graph is a time graph it tells us that during 2000 the Maximum\_Temperature is the least recorded according to the data. It is a scattered graph

In [49]: *# Bivariate Bar Plots*

```
df.groupby('Month')['Precipitation'].mean().plot(kind='bar',figsize=(20,10))
```

Out[49]: <AxesSubplot:xlabel='Month'>




The graph showing is bar graph and it tells us that in 3rd month the recorded precipitation is the highest.

```
In [50]: # Bivariate Bar Plots
```

```
df.groupby('Month')['Maximum_Temperature'].mean().plot(kind='bar',figsize=(20,10))
```

```
Out[50]: <AxesSubplot:xlabel='Month'>
```



The graph showing is bar graph and it tells us that in 12th month the recorded Maximum\_Temperature is the highest.




```
In [51]: # Correlation
import seaborn as sns
df.corr()

# Heatmap

plt.figure(figsize=(10,10))

sns.heatmap(df.corr(),annot=True)
```

Out[51]: <AxesSubplot:>



## **Splitting the dataset into training and testing**

```
In [52]: import datetime as dt
from sklearn.model_selection import train_test_split

year = pd.to_datetime(df.Date).dt.year
train_df = df[year == 2009]
val_df = df[year == 2013]
test_df = df[year == 2011]

train_df = train_df.drop(['Date'], axis=1)
val_df = val_df.drop(['Date'], axis=1)
test_df = test_df.drop(['Date'], axis=1)

train_df.info()
print(len(train_df), len(val_df), len(test_df))
```

<class 'pandas.core.frame.DataFrame'>  
DatetimeIndex: 274980 entries, 2009-01-01 to 2009-12-31  
Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	Precipitation	274980 non-null	float64
1	Maximum_Temperature	274980 non-null	float64
2	Minimum_Temperature	274980 non-null	float64
3	Relative_humidity	274980 non-null	float64
4	Wind_direction	274980 non-null	float64
5	Wind_speed	274980 non-null	float64
6	Cloudiness	274980 non-null	float64
7	Year	274980 non-null	int64
8	Month	274980 non-null	int64
9	Day	274980 non-null	int64
10	Day Name	274980 non-null	object

dtypes: float64(7), int64(3), object(1)  
memory usage: 25.2+ MB  
274980 265870 274561

## Making Target column

## Making RainOrNot as the target column and all other column as input columns

```
In [53]: def f(row):
    if row['Precipitation'] > 0:
        val = 1
    else:
        val = 0
    return val
```

```
In [54]: train_df['RainOrNot'] = train_df.apply(f, axis = 1)
test_df['RainOrNot'] = test_df['Precipitation'].map(lambda x: 1 if x > 0.0 else 0)
val_df['RainOrNot'] = val_df['Precipitation'].map(lambda x: 1 if x > 0.0 else 0)
```

```
In [73]: input_cols = list(train_df.columns)[1:7]
target_col = 'RainOrNot'
print('Input Features: ', input_cols)
print()
print('Target Feature: ', target_col)
```

Input Features: ['Maximum\_Temperature', 'Minimum\_Temperature', 'Relative\_humidity', 'Wind\_direction', 'Wind\_speed', 'Cloudiness']

Target Feature: RainOrNot

```
In [95]: train_inputs = train_df[input_cols].copy()
train_targets = train_df[target_col].copy()

val_inputs = val_df[input_cols].copy()
val_targets = val_df[target_col].copy()

test_inputs = test_df[input_cols].copy()
test_targets = test_df[target_col].copy()
```

## Implementing Logistic Regression

```
In [99]: from sklearn.linear_model import LogisticRegression
scores_dict={}
model = LogisticRegression(solver='liblinear')
model.fit(train_inputs, train_targets)

train_preds = model.predict(train_inputs)
train_probs = model.predict_proba(train_inputs)
```

```
In [100]: from sklearn.metrics import accuracy_score
accuracy_score(train_targets, train_preds)
```

Out[100]: 0.7866353916648483

```
In [92]: from sklearn.metrics import confusion_matrix

def predict_and_plot(inputs, targets, name=''):
    preds = model.predict(inputs)


    accuracy = accuracy_score(targets, preds)
    print("Accuracy: {:.2f}%".format(accuracy * 100))

    cf = confusion_matrix(targets, preds, normalize='true')
    plt.figure()
    sns.heatmap(cf, annot=True)
    plt.xlabel('Prediction')
    plt.ylabel('Target')
    plt.title('{} Confusion Matrix'.format(name));

    return preds, accuracy
```


```
In [93]: #Training Accuracy
import seaborn as sns
train_preds = predict_and_plot(train_inputs, train_targets)
```

Accuracy: 78.66%



```
In [121]: #Testing Accuracy  
test_accuracy_lr = predict_and_plot(test_inputs, test_targets)  
scores_dict['Logistic Regression']=0.7802
```

Accuracy: 78.02%



## Implementing Desicion Tree

In [124]:

```
from sklearn.tree import DecisionTreeClassifier
from sklearn import metrics

clf_entropy = DecisionTreeClassifier(criterion = "entropy", random_state = 100, n

clf_entropy.fit(train_inputs, train_targets)

y_pred = clf_entropy.predict(test_inputs)

test_accuracy_dt = metrics.precision_score(test_targets, y_pred)
print("Accuracy:",test_accuracy_dt)
print("Precision:",metrics.precision_score(test_targets, y_pred))
print("Recall:",metrics.recall_score(test_targets, y_pred))

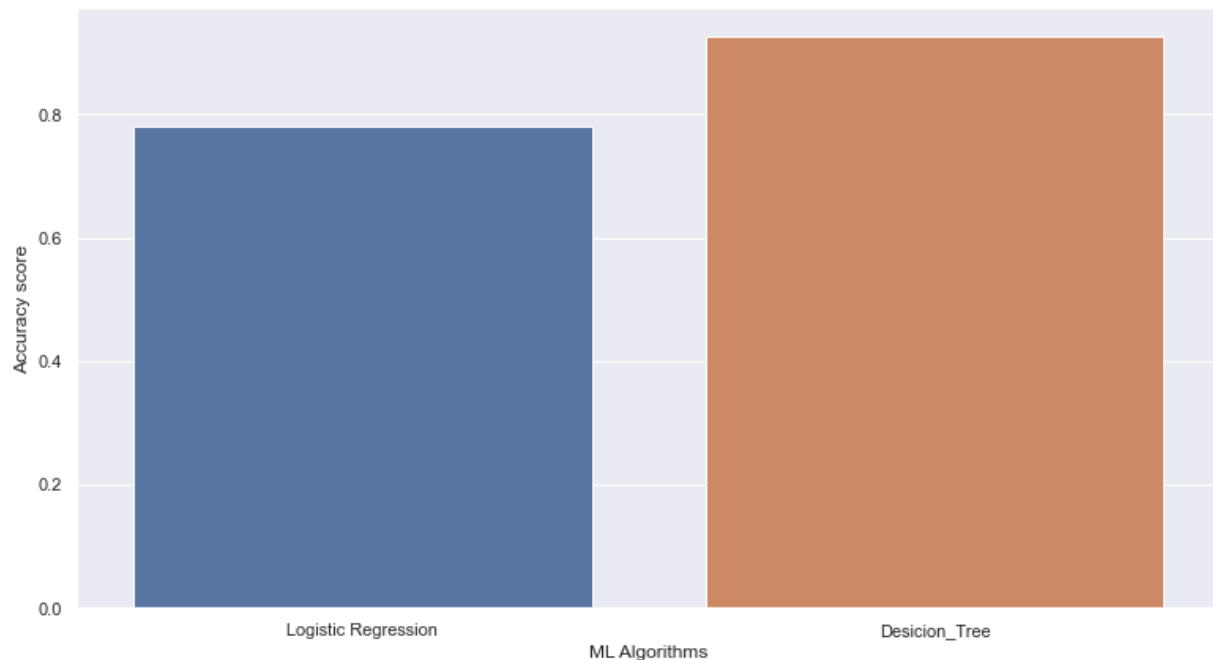
clf_entropy.fit(val_inputs, val_targets)
y_pred = clf_entropy.predict(val_inputs)
validation_accuracy_dt = metrics.precision_score(val_targets, y_pred)
scores_dict['Desicion_Tree']=test_accuracy_dt
```

Accuracy: 0.926561709228686  
Precision: 0.926561709228686  
Recall: 0.9480136590099084

```
In [123]: with sns.color_palette('muted'):
          algo_name = list(scores_dict.keys())
          scoress = list(scores_dict.values())

          sns.set(rc={'figure.figsize':(13,7)})
          plt.xlabel("ML Algorithms")
          plt.ylabel("Accuracy score")

          sns.barplot(algo_name,scoress)
```



**Conclusion: Among the 2 Algorithms Decision Tree has the highest accuracy.**

In [ ]: